

POLI 30 D: Political Inquiry
TA Sessions

Lab 07 | R Plots and R Data Analysis III

Before we start

Announcements:

- ▶ GitHub page:
<https://github.com/umbertomig/POLI30Dpublic>
- ▶ Piazza forum: The link in the slides needs to be fixed.
Check with instructors for an alternative link.

Before we start

Recap: In the Lab sessions, you learned:

- ▶ How to install R and R Studio on your computer.
- ▶ How to do basic and advanced operations with vectors and data frames.
- ▶ How to install packages and work with R Markdown.
- ▶ How to create plots and data viz.
- ▶ How to do data analysis.

Great job!

- ▶ Do you have any questions about these contents?

Plan for Lab 07

- Selecting variables
- Filtering cases in the dataset
- Grouping and summarizing data
- Visualizing correlations

Getting started

Getting started

- ▶ To get started, we need to load the datasets we will need in the lab.
- ▶ We also need to load the `tidyverse` package, which has all the R functions we use.

```
library(tidyverse)
```

Getting started - Education expenditure data

```
educexp <- read.csv("https://raw.githubusercontent.com/umberton  
head(educexp)
```

##	education	income	young	urban	states
## 1	189	2824	350.7	508	ME
## 2	169	3259	345.9	564	NH
## 3	230	3072	348.5	322	VT
## 4	168	3835	335.3	846	MA
## 5	180	3549	327.1	871	RI
## 6	193	4256	341.0	774	CT

Getting started - Chile survey data

```
chilesurv <- read.csv("https://raw.githubusercontent.com/umbert  
head(chilesurv)
```

##	statusquo	vote	voteYES
## 1	3.02460	Y	1
## 2	-3.88851	N	0
## 3	3.69216	Y	1
## 4	-3.09489	N	0
## 5	-3.31488	N	0
## 6	-3.14055	N	0

Getting started - Voting

```
voting <- read.csv("https://raw.githubusercontent.com/umbertomi  
head(voting)
```

##	birth	message	voted
## 1	1981	no	0
## 2	1959	no	1
## 3	1956	no	1
## 4	1939	yes	1
## 5	1968	no	0
## 6	1967	no	0

The tidyverse package

The tidyverse package

- ▶ The tidyverse package is an eco-system for data analytics.
- ▶ It has packages for processing, plot, and wrangle data.
- ▶ We have seen the package for plotting: `ggplot2`.
- ▶ But there are other important packages in the tidyverse constellation.
- ▶ And in the links we provide URLs for cheat sheets.

The tidyverse package

Other important packages:

- ▶ `forcats`: Package to deal with categorical variables.
- ▶ `stringr`: Package to deal with strings and texts in R.
- ▶ `purrr`: Package to work with functions and lists.
- ▶ `tidyr`: Package to clean up and reshape datasets.
- ▶ `dplyr`: Package to manipulate (wrangling) datasets.

We are going to do some `dplyr` today!

dplyr

- ▶ dplyr is the package on tidyverse for data wrangling.
- ▶ Data wrangling is a fancy way to say that you manipulate your dataset until it is in shape for analysis.
- ▶ Most of the time, we collect data from the internet, documents, or even old archives.
- ▶ We then need to do some *wrangling* to ensure it is ready to extract the statistics of interest.

dplyr

dplyr is based on a bunch of verbs that do what we intuitively think each of them should do:

- ▶ Variable/column:
 - ▶ `select`: Select variables
 - ▶ `rename`: Rename variables
 - ▶ `mutate`: Mutate variables
- ▶ Observation/row:
 - ▶ `filter`: Filter dataset
 - ▶ `arrange`: Arrange observations
- ▶ Get summaries:
 - ▶ `summarize`: Summarize the dataset
 - ▶ `group_by`: Group summary by a variable

Selecting Variables (column operations)

Selecting Variables (column)

Suppose we want to select only the numeric variables in the educexp dataset. We do the following:

```
educ2 <- select(educexp, education, income, young, urban)
head(educ2, 2)
```

```
##      education income young urban
## 1          189   2824 350.7   508
## 2          169   3259 345.9   564
```

► The syntax is:

```
newdat <- select(olddat, v1, ..., vn)
```


Selecting Variables (column)

Or a more compact notation for this uses *pipes*:

```
educ3 <- educexp %>%  
  select(education, income, young, urban)  
head(educ3, 2)  
##      education income young urban  
## 1          189   2824 350.7   508  
## 2          169   3259 345.9   564
```

- Note that the *pipe* operator (`%>%`) works like a composite function: takes what is in the left and passes through to the right.
- The first thing that the `select` is looking for is the dataset. The *pipe* passes the dataset, so you do not need to mention it.

Selecting Variables (column)

And the methods we can apply in the select are the following:

Method	Effect
v1, v2, v3 (etc)	Select given variables
starts_with('xyz')	Select starting with xyz
ends_with('xyz')	Select ending with xyz
contains('xyz')	Select variables that have xyz in their names
vk:vn	All variables between vk and vn
-(vk:vn)	All but the variables between vk and vn

We can even rename variables using select.

Selecting Variables (column)

Example 1 (same selection as before):

```
educ4 <- educexp %>% select(-states)
```

```
head(educ4, 3)
```

```
##      education income young urban
```

```
## 1          189    2824 350.7   508
```

```
## 2          169    3259 345.9   564
```

```
## 3          230    3072 348.5   322
```

Example 2 (in the Chile dataset now):

```
chile2 <- chilesurv %>% select(starts_with('vote'))
```

```
head(chile2, 3)
```

```
##      vote voteYES
```

```
## 1      Y         1
```

```
## 2      N         0
```

```
## 3      Y         1
```

Rename

We can change the names of the variables using `rename`. The basic syntax is:

```
newdat <- olddat %>%  
  rename(newname1 = oldname1,  
         newname2 = oldname2,  
         etc)
```

Example:

```
chile3 <- chilesurv %>% rename(votebinary = voteYES)  
head(chile3, 2)  
##      statusquo vote votebinary  
## 1      3.02460    Y            1  
## 2     -3.88851    N            0
```

Pipeing like a boss

You can *pipe* as many commands as you want. Example:

```
chile4 <- chilesurv %>%  
  select(starts_with('vote')) %>%  
  rename(votebinary = voteYES)  
head(chile4, 2)
```

##	vote	votebinary
## 1	Y	1
## 2	N	0

► Note the *chain* pipe!

Filtering Rows

Filtering Rows

- ▶ Filtering the data means subset the dataset based on a logical condition.
 - ▶ Example: In the voting pressure experiment, we may want to study the effect of messaging young people.
 - ▶ Reason: Younger people may not be as sensitive to peer pressure.

```
voting2 <- voting %>% filter(birth >= 1975)  
head(voting2, 4)
```

##	<i>birth</i>	<i>message</i>	<i>voted</i>
## 1	1981	no	0
## 2	1983	no	0
## 3	1979	no	0
## 4	1985	yes	1

Filtering Rows

► Syntax:

```
newdat <- filter(olddat, condition1, condition2, etc)
or
newdat <- olddat %>% filter(condition1, condition2,
etc)
```

- Note that you may or may not use the pipe operator. But we can all agree that it makes the code much cleaner.

Filtering rows

- The filter operators are the following:

Operator	Meaning
< or <=	Smaller than or smaller than or equal
> or >=	Greater than or greater than or equal
==	Equal
!=	Different
!	Negation (turns a TRUE into a FALSE)
	Or
&	And

Arrange

- We can sort the dataset by variable content according to the needs of our analysis.

```
newdat <- olddat %>% arrange(var1, desc(var2), etc)
```

- The `desc()` command sorts descending (highest to the lowest). The default is to sort ascending.

Arrange

- Example: Sorting by education expenditure (high to low):

```
educexp %>% arrange(desc(education)) %>% head()
```

##	<i>education</i>	<i>income</i>	<i>young</i>	<i>urban</i>	<i>states</i>
## 1	372	4146	439.7	484	AK
## 2	273	3968	348.4	909	CA
## 3	262	3341	365.4	664	MN
## 4	261	4151	326.2	856	NY
## 5	248	3795	375.9	722	DE
## 6	247	3742	364.1	766	MD

Mutating (altering the content)

Mutating

- ▶ We frequently need to operate with our variables. For example, suppose that you want to apply a log function to `income`.
- ▶ `log` has the advantage of transforming large changes into smaller ones.
- ▶ It also has the advantage of interpreting regression models. Check [out this UofV Library entry](#).
- ▶ Syntax:

```
newdat <- olddat %>% mutate(vnew = calcs(vold), etc)
```

Mutating

Example: Log of education and the sum of variables:

```
educ5 <- educexp %>%  
  mutate(logeduc = log(education),  
         sum_iyu = income + young + urban)  
head(educ5, 4)
```

##	education	income	young	urban	states	logeduc	sum_iyu
## 1	189	2824	350.7	508	ME	5.241747	3682.7
## 2	169	3259	345.9	564	NH	5.129899	4168.9
## 3	230	3072	348.5	322	VT	5.438079	3742.5
## 4	168	3835	335.3	846	MA	5.123964	5016.3

Your turn: Adapt this code to arrange the dataset in terms of logeduc.

Grouping(-by) and Summarizing

Summarize

- ▶ The `summarize` function is helpful when computing summary statistics in the whole sample.
- ▶ Syntax:

```
olddat %>% summarize(stat1 = calcs(vars1), stat2 = calcs(vars
```

- ▶ Note that we only save the results if we need them! The idea is that those summary statistics you compute to check, not save.

Summarize

- Example: Suppose we want to compute the mean, standard deviation, and the number of observations for education and income.

```
educexp %>%  
  summarize(avgeduc = mean(education),  
            stdeveduc = sd(education), neduc = n(),  
            avginc = mean(income),  
            stdevinc = sd(income), ninc = n())  
##      avgeduc stdeveduc neduc      avginc stdevinc ninc  
## 1 196.3137  46.45449    51 3225.294  560.026    51
```

- Besides the number of observations, these are pretty informative statistics.

Group-by and Summarize

- ▶ But we may want the summary by a group. For example, to compute the *differences-in-means* estimator, we need:
 - ▶ The average in the treatment group
 - ▶ The average in the control group.
- ▶ Group-by can help us here:
 - ▶ First, we group our results by the treatment status.
 - ▶ Then, we summarize.
- ▶ It will create one summary for each group.
- ▶ Syntax:

```
dat %>% group_by(groupvar) %>%  
  summarize(stat1 = calcs(vars1), etc)
```

Group-by and Summarize

► Example:

```
voting %>% group_by(message) %>%  
  summarize(avgvote = mean(voted) * 100, ncases = n())  
## # A tibble: 2 x 3  
##   message avgvote ncases  
##   <chr>      <dbl>   <int>  
## 1 no          29.7  191243  
## 2 yes         37.8   38201
```

- And the differences-in-means estimator is then straightforward to compute.
- Note that we also saved the number of cases in each treatment status.
 - It makes more sense: Number of control and treatment groups cases.

Group-by and Summarize

Your turn: Adapt this example here to compute the percentage of voters in the treatment and the control group, filtering by younger voters (born on or after 1975).

► Example:

```
voting %>%  
  group_by(message) %>%  
  summarize(avgvote = mean(voted) * 100,  
            ncases = n())
```

Question: Does the treatment effect increase or decrease?
Why do you think that is the case?

Detour: **tibble** versus **data.frame**

- ▶ When using `dplyr` functions such as the ones we have used in here, `tidyverse` sometimes saves the results in `tibbles`.
- ▶ `tibbles` are pretty much the same as `data.frames`.
- ▶ To go back to `data.frame`, you should do the following:
`dfdat <- data.frame(tibbledat)`.
- ▶ Or you can add a `%>% data.frame()` to the end of your piping. For this class, it does not matter which one you use.
- ▶ If you want to check about `tibbles`, [here](#) is a good source.

Plotting Correlations

ggcorrplot and GGally

- ▶ These packages are helpful to add some more spice to our plots.
- ▶ For example, ggcorrplot is good for plotting correlations of multiple variables.
 - ▶ The command `cor` computes two-by-two correlations for all the combinations of variables when more than two vars.

▶ Installing packages: Run this code in your console:

```
install.packages('ggcorrplot')  
install.packages('GGally')
```

- ▶ **Warning:** You should *never* install packages on R Markdown!

ggcorrplot and GGally

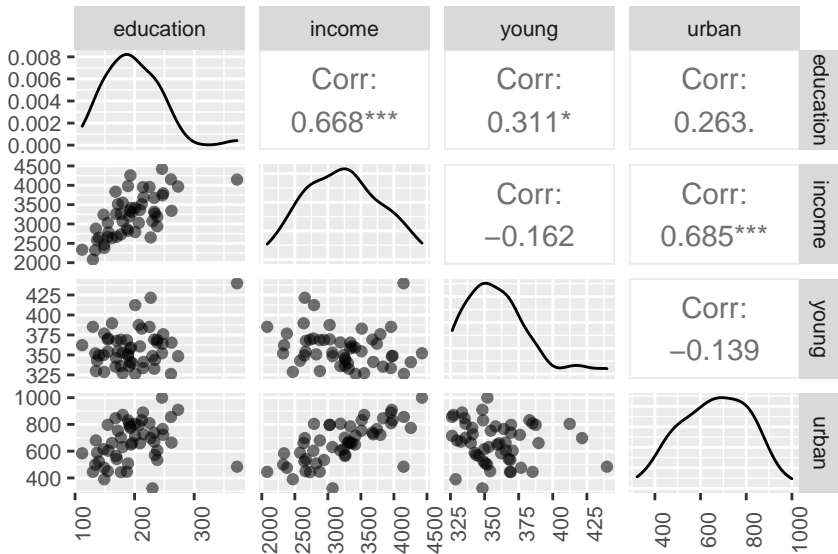
- Now, to load both packages:

```
library('ggcorrplot')  
library('GGally')
```

```
## Registered S3 method overwritten by 'GGally':  
##   method from  
##   +.gg      ggplot2
```

- And we now have all the functionalities of these packages at our disposal!


```
ggpairs(educexp %>% select(-states), aes(alpha = 0.4)) +  
  theme(text = element_text(size = 10),  
        axis.text.x = element_text(angle = 90) )
```



Correlations

- And the correlation between these variables is:

```
educexp %>% select(-states) %>% cor()
```

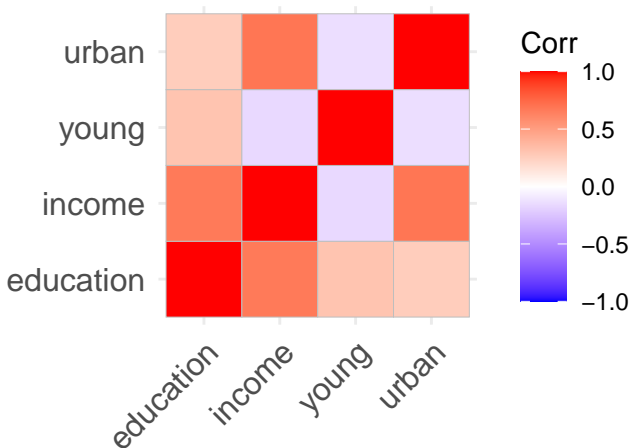
```
##           education      income      young      urban
## education 1.0000000 0.6675773 0.3114855 0.2633238
## income    0.6675773 1.0000000 -0.1623600 0.6854580
## young     0.3114855 -0.1623600 1.0000000 -0.1386334
## urban     0.2633238 0.6854580 -0.1386334 1.0000000
```

- The correlation between the variable and itself is always one.
- The correlation is a **symmetric matrix**:
 - The correlation between income and education is the same as the correlation between education and income.

Correlations

Using squares:

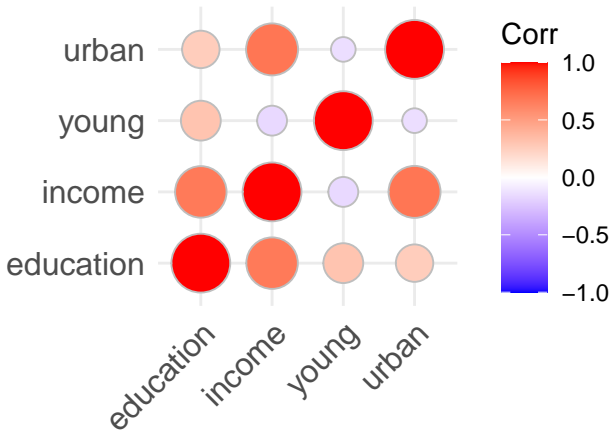
```
ggcorrplot(educexp %>% select(-states) %>% cor())
```



Correlations

And using circles:

```
ggcorrplot(educexp %>% select(-states) %>% cor(),  
            method = 'circle')
```



ggcorrplot and GGally

- ▶ Note that both commands use **only numeric variables**.
- ▶ Another thing to note is that you can do correlation plots using GGally.
 - ▶ Function name: `ggcorr`. The syntax is different, but it also produces a nice plot.
- ▶ Syntax for `ggpairs`:

```
dat %>% select(numericvars) %>% ggpairs()
```

- ▶ Syntax for `ggcorrplot`:

```
dat %>% select(numericvars) %>%  
  cor() %>% ggcorrplot()
```

Today's Lab

- Selecting variables
- Filtering cases in the dataset
- Grouping and summarizing data
- Visualizing correlations

Next Lab

- More data wrangling
- One more plot
- Dealing with missing data
- Lots of analysis

Questions?

See you in the next lab!